



# Credit Card Delinquency Prediction using ML

By Atul Pai

# What is Credit Delinquency?

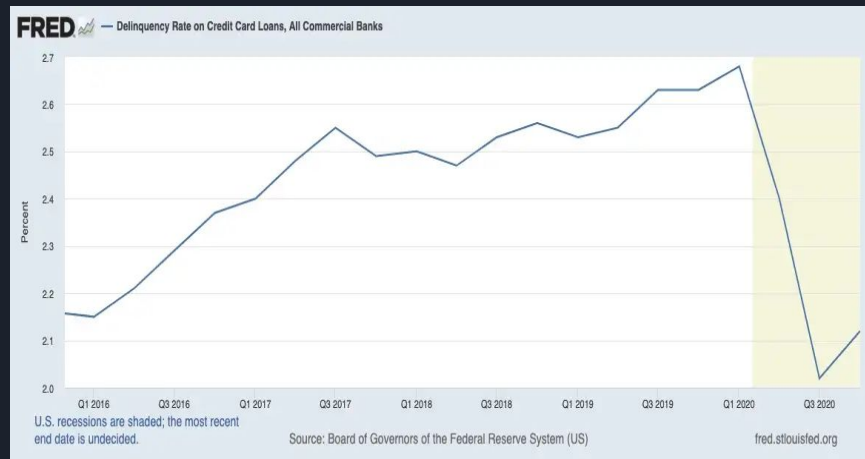
## Delinquency vs Default:


### Delinquency:

In general is a slightly mild term where a borrower is behind on payments. (usually 60 days)

### Default :

Once a person or institution is delinquent for over a specified period (usually 9 months).





# Why does it matter to Financial Institutions? What do they do?

## Reasons to Monitor:

- Macro Factors : Loan delinquencies which eventually lead to default can cause Global Financial Crisis.
- Reputation: Analysts measure a bank based on their Delinquency rates, affects their appearance in the eyes of their shareholders.
- Operational Costs : Bear collection and loss of capital costs .

## Factors generally used to predict:

- Credit Score/FICO Score
- Interest Rate
- Debt-to-income ratio of the borrower
- Number of days with a credit line
- Borrower's revolving balance
- Utilization rate of the borrower's revolving line
- Number of times the borrower had not paid in full or gone 30+ days past the due date on payment in the last two years.



# Dataset :

## Kaggle Competition : Give Me Some Credit

| Variable Name                        | Description  | Type       |
|--------------------------------------|--|------------|
| SeriousDlqin2yrs                     | Person experienced 90 days past due delinquency or worse   | Y/N        |
| RevolvingUtilizationOfUnsecuredLines | Total balance on credit cards and personal lines of credit except real estate and no installment debt like car loans divided by the sum of credit limits | percentage |
| age                                  | Age of borrower in years   | integer    |
| NumberOfTime30-59DaysPastDueNotWorse | Number of times borrower has been 30-59 days past due but no worse in the last 2 years.  | integer    |
| DebtRatio                            | Monthly debt payments, alimony, living costs divided by monthly gross income   | percentage |
| MonthlyIncome                        | Monthly income   | real       |
| NumberOfOpenCreditLinesAndLoans      | Number of Open loans (installment like car loan or mortgage) and Lines of credit (e.g. credit cards)   | integer    |
| NumberOfTimes90DaysLate              | Number of times borrower has been 90 days or more past due.  | integer    |
| NumberRealEstateLoansOrLines         | Number of mortgage and real estate loans including home equity lines of credit   | integer    |
| NumberOfTime60-89DaysPastDueNotWorse | Number of times borrower has been 60-89 days past due but no worse in the last 2 years.  | integer    |
| NumberOfDependents                   | Number of dependents in family excluding themselves (spouse, children etc.)  | integer    |

Problem:

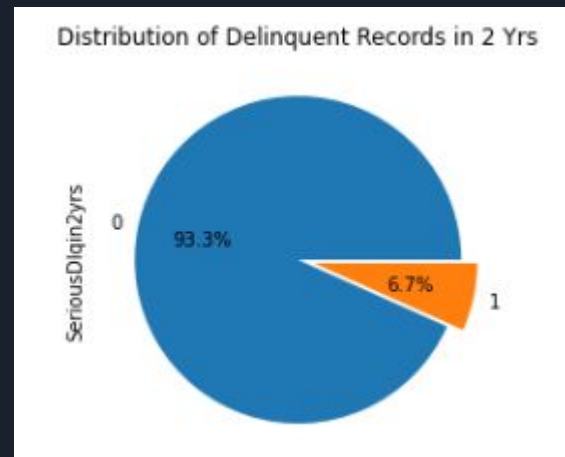
Based on the above features predict if an applicant will be delinquent in the next 2 years.

Number of Records : 150000

# Exploratory Data Analysis:

| #  | Column                               | Non-Null Count  | Dtype   |
|----|--------------------------------------|-----------------|---------|
| 0  | Unnamed: 0                           | 150000 non-null | int64   |
| 1  | SeriousDlqin2yrs                     | 150000 non-null | int64   |
| 2  | RevolvingUtilizationOfUnsecuredLines | 150000 non-null | float64 |
| 3  | age                                  | 150000 non-null | int64   |
| 4  | NumberOfTime30-59DaysPastDueNotWorse | 150000 non-null | int64   |
| 5  | DebtRatio                            | 150000 non-null | float64 |
| 6  | MonthlyIncome                        | 120269 non-null | float64 |
| 7  | NumberOfOpenCreditLinesAndLoans      | 150000 non-null | int64   |
| 8  | NumberOfTimes90DaysLate              | 150000 non-null | int64   |
| 9  | NumberRealEstateLoansOrLines         | 150000 non-null | int64   |
| 10 | NumberOfTime60-89DaysPastDueNotWorse | 150000 non-null | int64   |
| 11 | NumberOfDependents                   | 146076 non-null | float64 |

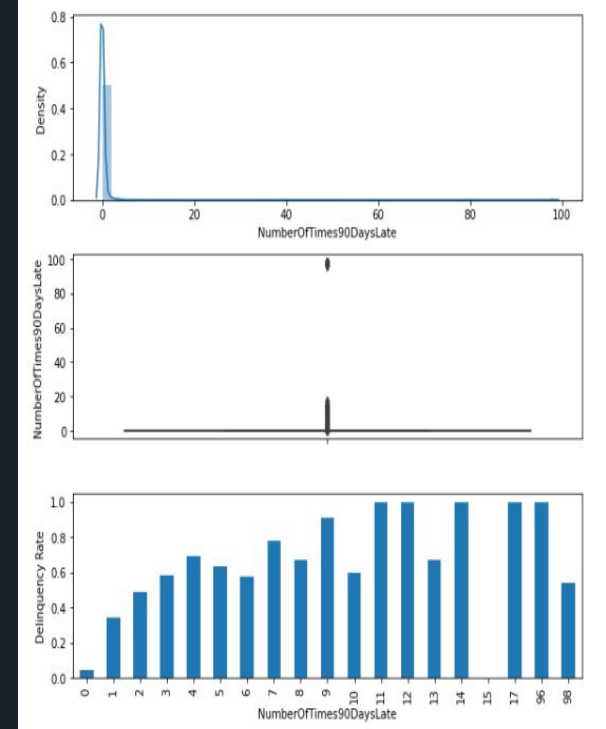
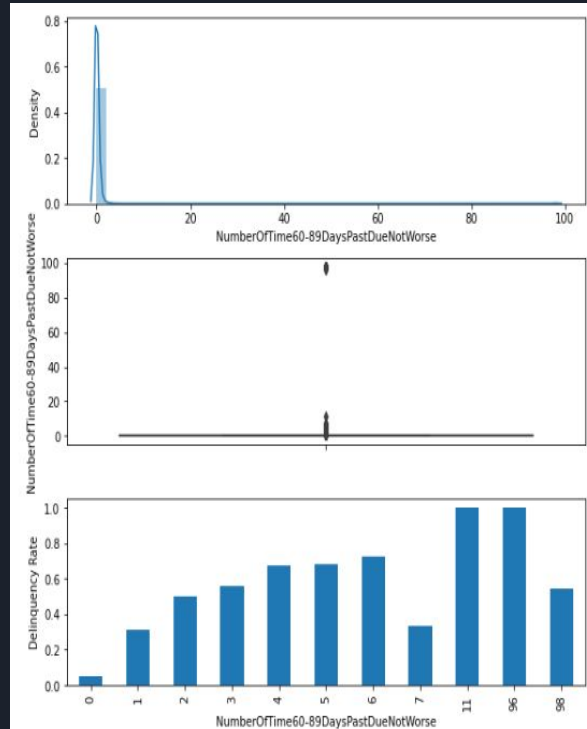
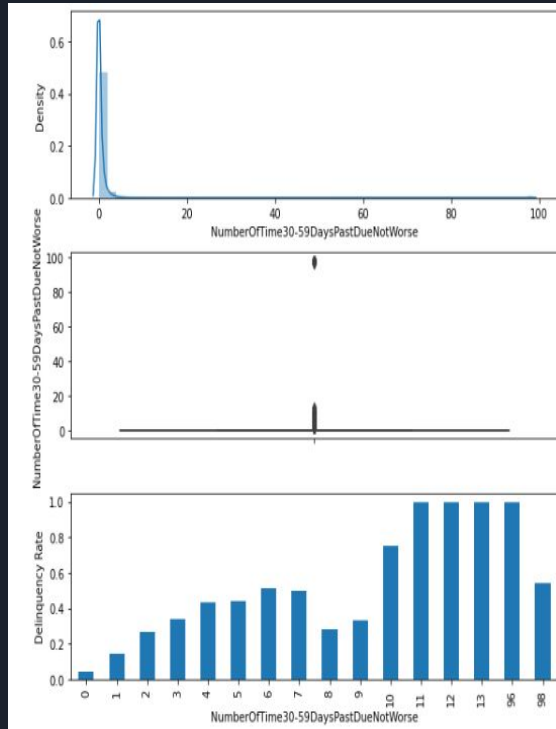
- Highly Imbalanced Dataset
- Null Values



# Exploratory Data Analysis: Univariate Analysis

See how the target variable changes with each of the explanatory variables:

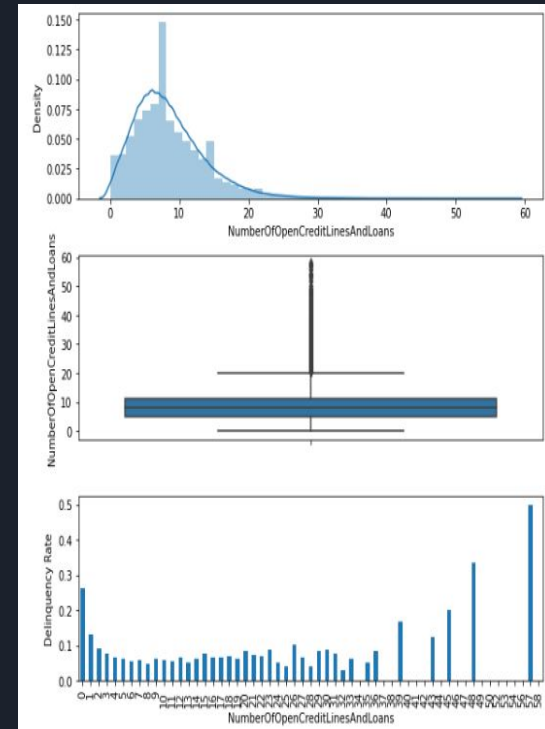
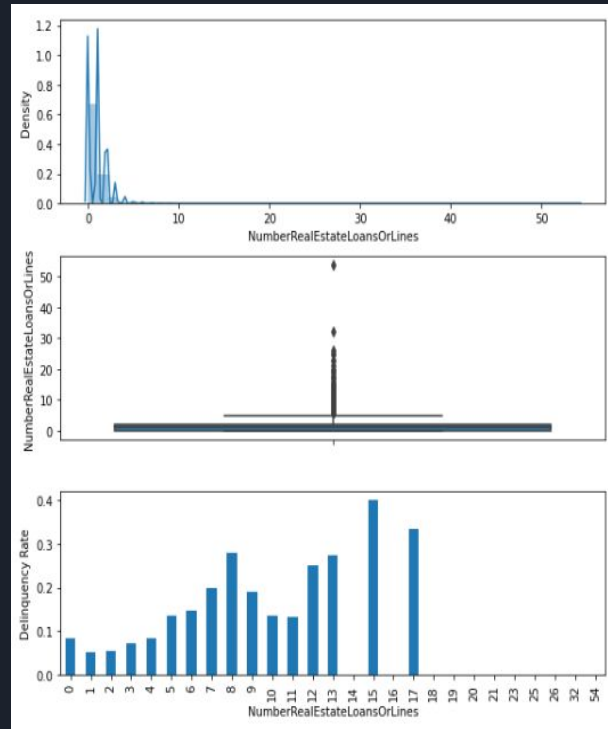
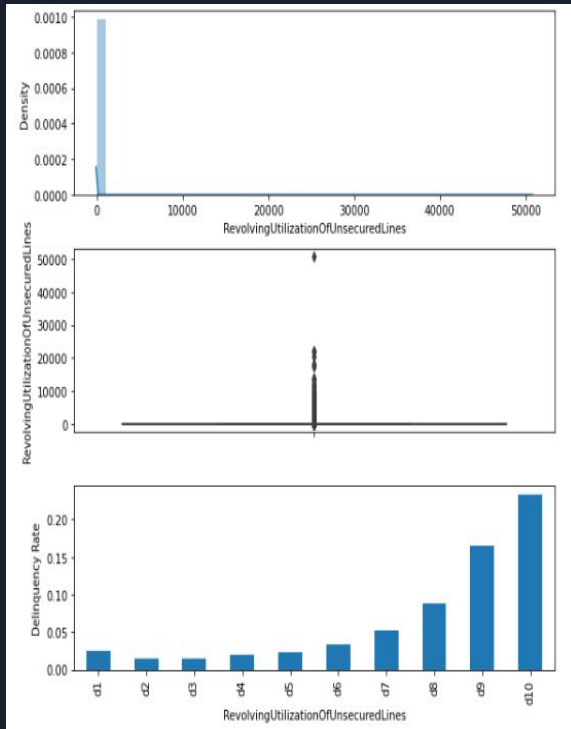
Frequency Features:



# Exploratory Data Analysis: Univariate Analysis

See how the target variable changes with each of the explanatory variables:

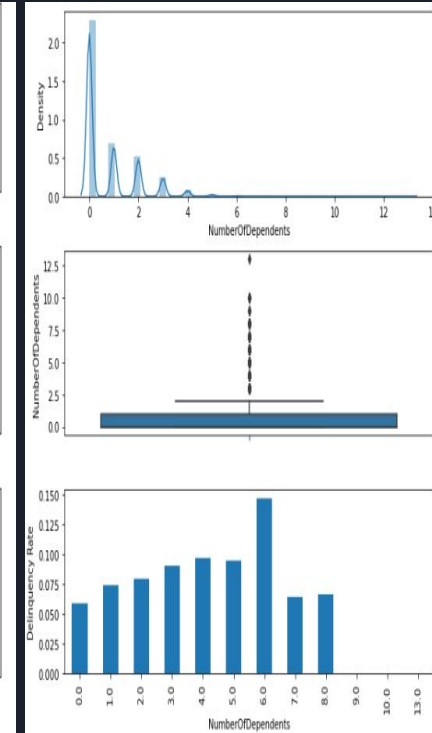
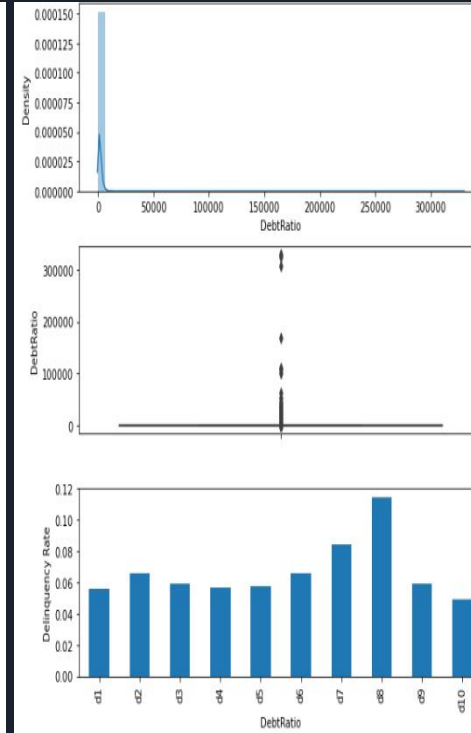
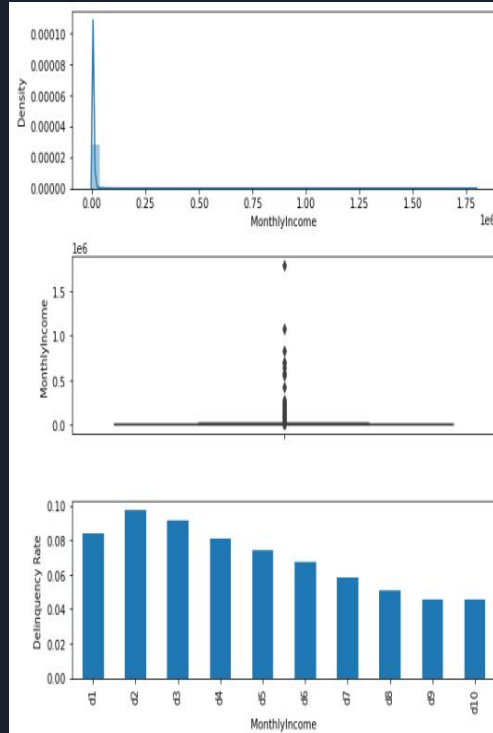
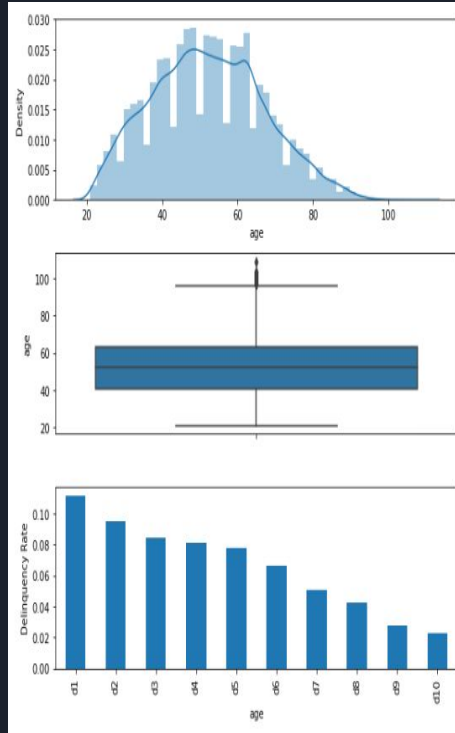
Revolving Utilization, Number of Real Estate Loans, Number of Credit Lines



# Exploratory Data Analysis: Univariate Analysis

See how the target variable changes with each of the explanatory variables:

Age, Monthly Income, Debt Ratio, Number of Dependents





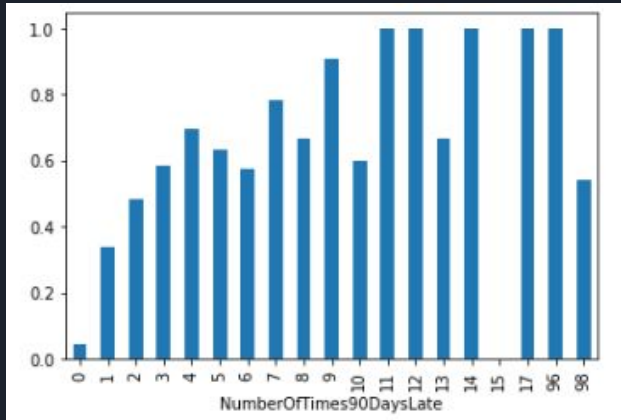
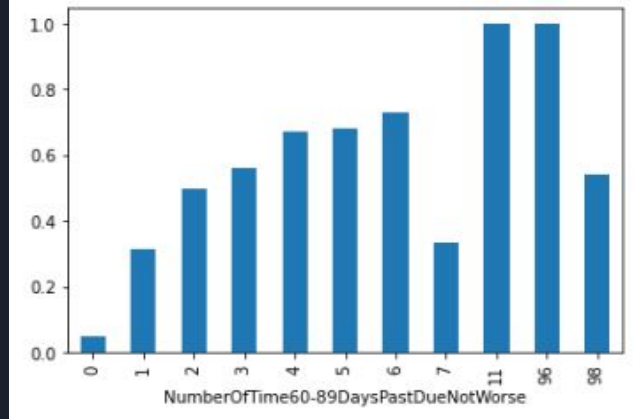
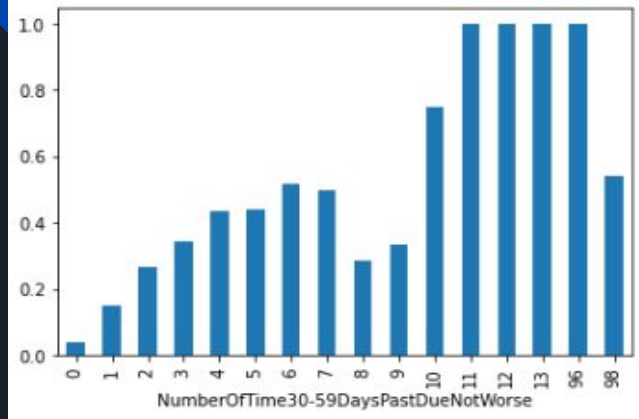
# Exploratory Data Analysis: Univariate Analysis

See how the target variable changes with each of the explanatory variables:

## Summary

| Variable                               | Trend                              | Range of Delinquency Rate | Rank |
|--|------------------------------------|---------------------------|------|
| 'NumberOfTimes90DaysLate'              | Increases with number of days      | 2% - 100%                 | 1    |
| 'NumberOfTime60-89DaysPastDueNotWorse' | Increases with number of days      | 2% - 100%                 | 2    |
| 'NumberOfTime30-59DaysPastDueNotWorse' | Increases with number of days      | 2% - 100%                 | 3    |
| 'RevolvingUtilizationOfUnsecuredLines' | Increases with increase in deciles | 2.5% - 25%                | 4    |
| 'NumberRealEstateLoansOrLines'         | Increases with number of loans     | 6%-30%                    | 5    |
| 'NumberOfOpenCreditLinesAndLoans'      | Highest for 0                      | 25% - 6%                  | 6    |
| 'age'                                  | Decrease with decrease in deciles  | 12% - 3%                  | 7    |
| 'MonthlyIncome'                        | Decrease with decrease in deciles  | 4%-8%                     | 8    |
| 'DebtRatio'                            | No strong trend in particular      | ~ 6%                      | 9    |
| 'NumberOfDependents'                   | Remains steady                     | ~ 6%                      | 10   |

# Exploratory Data Analysis: Cleaning



- 96,98 labeled values drop
- Income: replace null with Median
- Number of Dependents : replace null with 0 median
- No significant correlation between features



## Metric of consideration:

- Chose Class 1 recall as the metric to optimize:
  - Aim to catch as many delinquencies as possible
  - A high recall for class 1 would mean low False Negatives for Delinquent accounts.
  - Drawback is that it could lead to high False positives, which would mean additional work.

## Engineered Features:

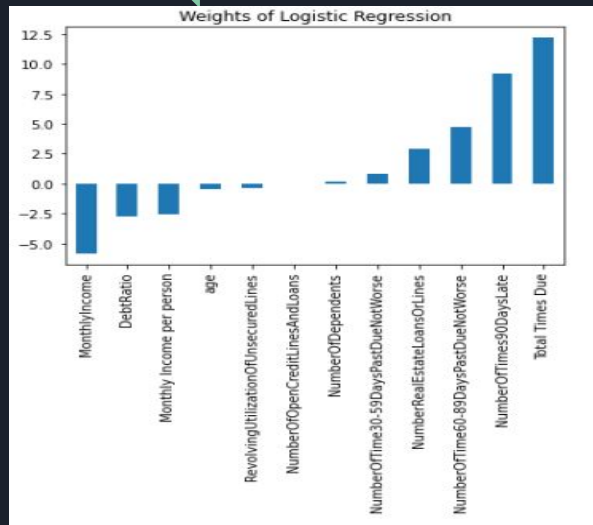
- Total Times Due =  $N_{30-60} + N_{60-90} + N_{90}$
- Monthly Income/ Person

Baseline: XGBoost with no oversampling

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.94      | 0.99   | 0.97     | 30922   |
| 1            | 0.56      | 0.17   | 0.27     | 2184    |
| accuracy     |           |        | 0.94     | 33106   |
| macro avg    | 0.75      | 0.58   | 0.62     | 33106   |
| weighted avg | 0.92      | 0.94   | 0.92     | 33106   |

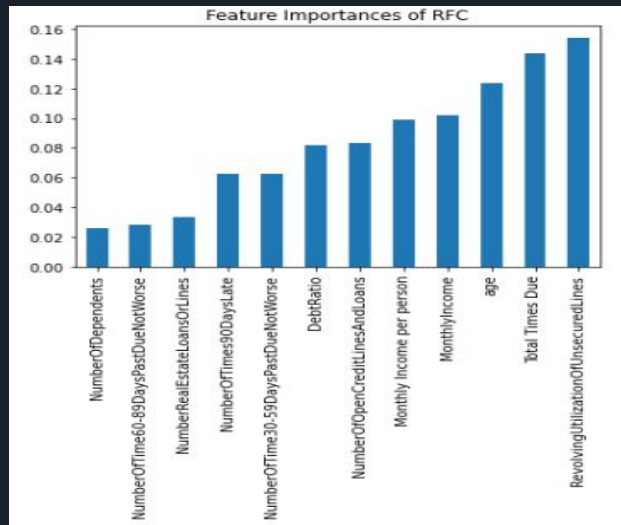
# Modeling - Oversampling:

Model 1: Logistic Regression with Oversampling



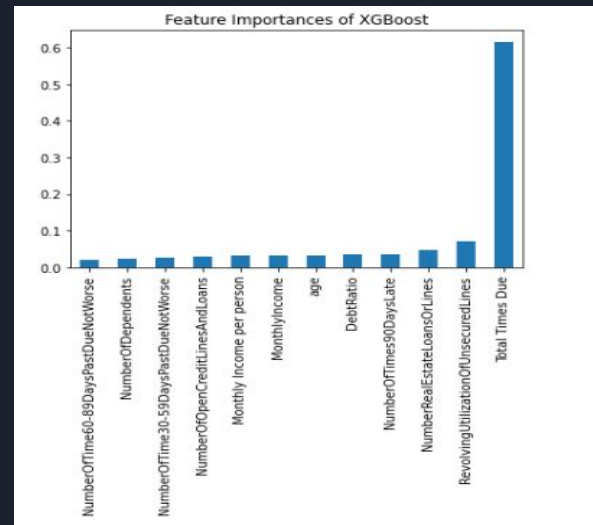
|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.97      | 0.88   | 0.92     | 30922   |
| 1            | 0.26      | 0.62   | 0.37     | 2184    |
| accuracy     |           |        | 0.86     | 33106   |
| macro avg    | 0.62      | 0.75   | 0.64     | 33106   |
| weighted avg | 0.92      | 0.86   | 0.88     | 33106   |

Model 2: Random Forest with Oversampling



|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.95      | 0.98   | 0.96     | 30922   |
| 1            | 0.45      | 0.24   | 0.32     | 2184    |
| accuracy     |           |        | 0.93     | 33106   |
| macro avg    | 0.70      | 0.61   | 0.64     | 33106   |
| weighted avg | 0.92      | 0.93   | 0.92     | 33106   |

Model 3: XGBoost with Oversampling

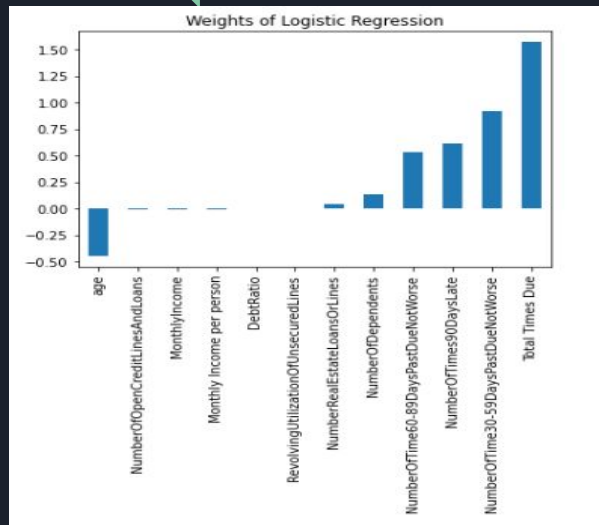


|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.97      | 0.85   | 0.91     | 30922   |
| 1            | 0.24      | 0.66   | 0.36     | 2184    |
| accuracy     |           |        | 0.84     | 33106   |
| macro avg    | 0.61      | 0.76   | 0.63     | 33106   |
| weighted avg | 0.92      | 0.84   | 0.87     | 33106   |

# Modeling - Oversampling and Tuning:

## Model 1: Logistic Regression with Oversampling

'C': 0.000774263682681127, penalty': 'l2'

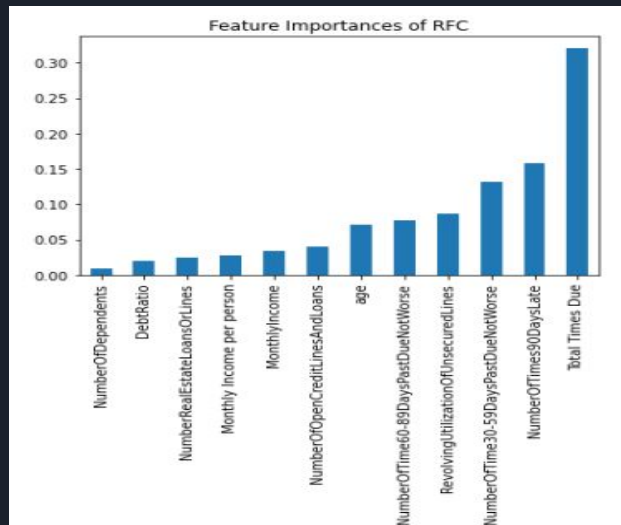


|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.97      | 0.66   | 0.79     | 30922   |
| 1            | 0.12      | 0.66   | 0.21     | 2184    |
| accuracy     |           |        | 0.66     | 33106   |
| macro avg    | 0.54      | 0.66   | 0.50     | 33106   |
| weighted avg | 0.91      | 0.66   | 0.75     | 33106   |

## Model 2: Random Forest with Oversampling

n\_estimators': 50,min\_samples\_split': 10

10,max\_depth': 10,min\_samples\_leaf': 4

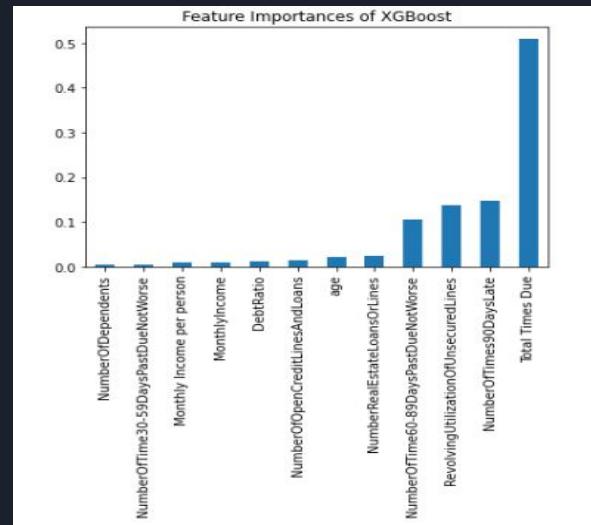


|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.97      | 0.87   | 0.92     | 30922   |
| 1            | 0.26      | 0.65   | 0.37     | 2184    |
| accuracy     |           |        | 0.86     | 33106   |
| macro avg    | 0.62      | 0.76   | 0.65     | 33106   |
| weighted avg | 0.93      | 0.86   | 0.88     | 33106   |

## Model 3: XGBoost with Oversampling

subsample: 0.8,max\_depth': 2,gamma': 0

colsample\_bytree': 0.8



|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.98      | 0.80   | 0.88     | 30922   |
| 1            | 0.21      | 0.76   | 0.33     | 2184    |
| accuracy     |           |        | 0.80     | 33106   |
| macro avg    | 0.60      | 0.78   | 0.61     | 33106   |
| weighted avg | 0.93      | 0.80   | 0.84     | 33106   |

# Modeling - Stack LR, RFC Vs XGBoost:

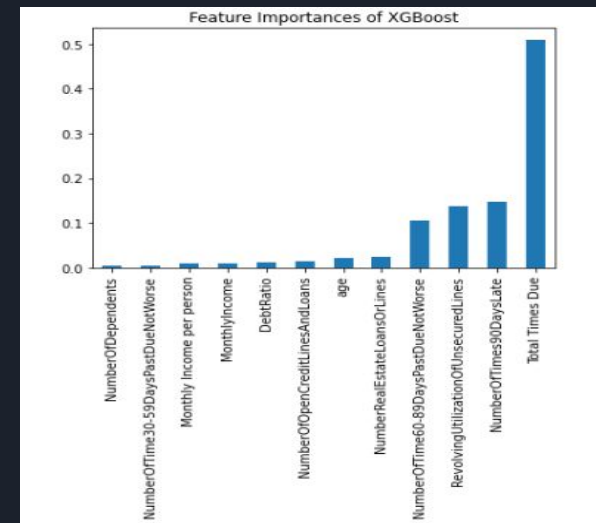
Model 1: Stacked LR, RFC

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.97      | 0.85   | 0.91     | 30922   |
| 1            | 0.25      | 0.68   | 0.36     | 2184    |
| accuracy     |           |        | 0.84     | 33106   |
| macro avg    | 0.61      | 0.77   | 0.64     | 33106   |
| weighted avg | 0.93      | 0.84   | 0.87     | 33106   |

Model 2: XGBoost with  
Oversampling

subsample: 0.8,max\_depth': 2,gamma': 0

colsample\_bytree': 0.8



|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.98      | 0.80   | 0.88     | 30922   |
| 1            | 0.21      | 0.76   | 0.33     | 2184    |
| accuracy     |           |        | 0.80     | 33106   |
| macro avg    | 0.60      | 0.78   | 0.61     | 33106   |
| weighted avg | 0.93      | 0.80   | 0.84     | 33106   |

# Predictions with XGBoost:

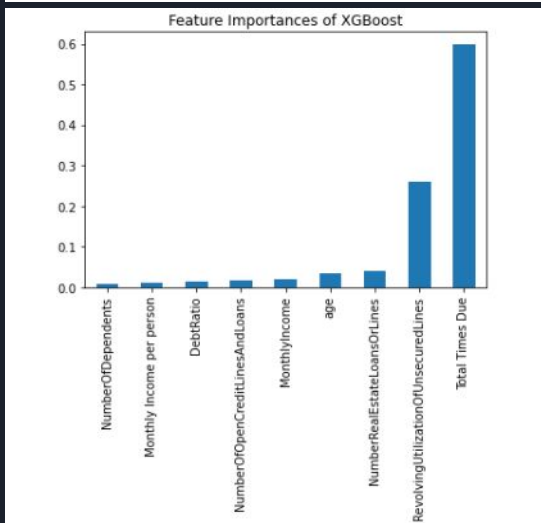
Final: XGBoost with

Oversampling

subsample: 0.8,max\_depth':

2,gamma': 0

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.98      | 0.79   | 0.87     | 46151   |
| 1            | 0.21      | 0.78   | 0.33     | 3260    |
| accuracy     |           |        | 0.79     | 49411   |
| macro avg    | 0.59      | 0.78   | 0.60     | 49411   |
| weighted avg | 0.93      | 0.79   | 0.84     | 49411   |

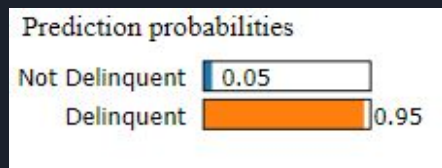
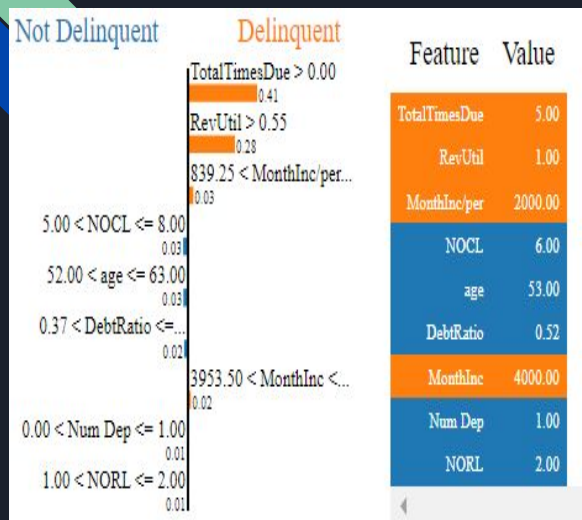


Interpretation:

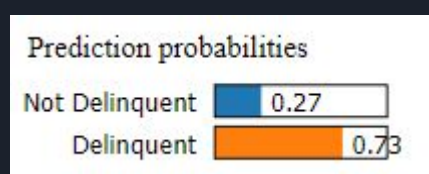
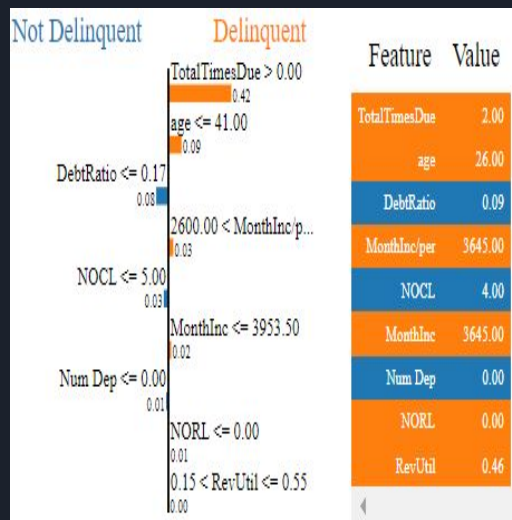
- The recall score for Class 1 is 78%. It means out of all the delinquent accounts, model catches 78% of them.
- The downside is out of all the predictions the model says is delinquent, 21% are actually delinquent.
- The model flags about 25% of the records as delinquent, analysing only these will help weed out 78% delinquent accounts.



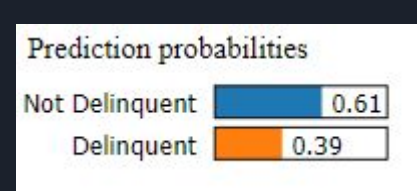
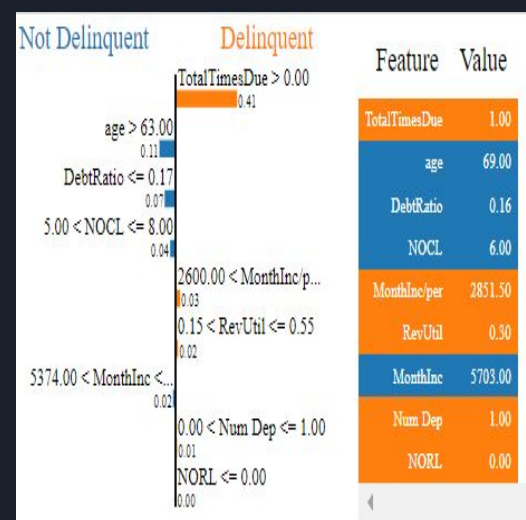
# Local Interpretability using LIME:



Correctly Classified: Delinquent predicted delinquent



Wrongly Classified: Non Delinquent predicted delinquent



Wrongly Classified: Delinquent predicted Non delinquent



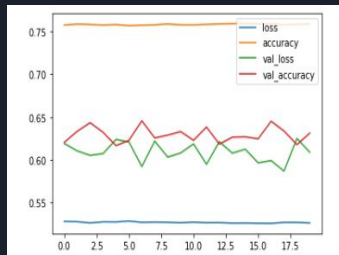
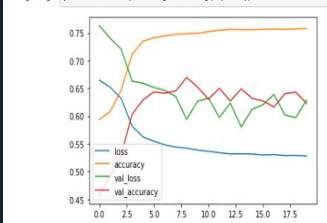
# Deep Learning: Used all features including the feature engineered ones with oversampler.

1) 8 nodes in the first layer and 8 in the second

```
model = Sequential()
model.add(Dense(8, activation='relu', input_shape=(n_inputs,)))
model.add(Dropout(0.25))
model.add(Dense(8, activation='relu'))
model.add(Dropout(0.25))
model.add(Dense(1, activation='sigmoid'))

model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

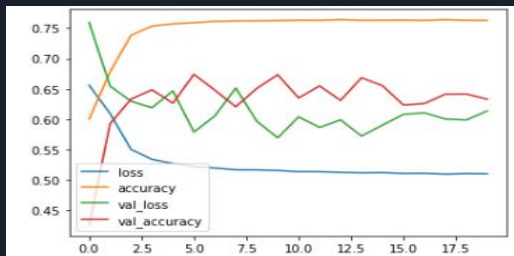
In [335]: pd.DataFrame(history.history).plot()



2) 16 nodes in the first layer and 16 in the second with

```
n_inputs = X_train_r.shape[1]
model = Sequential()
model.add(Dense(16, activation='relu', input_shape=(n_inputs,)))
model.add(Dropout(0.25))
model.add(Dense(16, activation='relu'))
model.add(Dropout(0.25))
model.add(Dense(1, activation='sigmoid'))

model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
history = model.fit(X_train_r, y_train_r, epochs=20, validation_split=0.1, batch_size=256)
```



|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.97      | 0.80   | 0.88     | 46151   |
| 1            | 0.20      | 0.70   | 0.31     | 3260    |
| accuracy     |           |        | 0.79     | 49411   |
| macro avg    | 0.59      | 0.75   | 0.60     | 49411   |
| weighted avg | 0.92      | 0.79   | 0.84     | 49411   |

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.97      | 0.79   | 0.87     | 46151   |
| 1            | 0.19      | 0.71   | 0.30     | 3260    |
| accuracy     |           |        | 0.78     | 49411   |
| macro avg    | 0.58      | 0.75   | 0.59     | 49411   |
| weighted avg | 0.92      | 0.78   | 0.83     | 49411   |

